



Core Idea

"There's no perfect answer — only trade-offs."



in @ankitrathi

Designing Data Systems

It's All About Trade-offs

⌚ Operational vs Analytical Systems

OLTP (Operational)

- Frequent writes
- Small queries
- Real-time transactions
- Use: Banking apps, Order systems



OLAP (Analytical)

- Complex reads
- Historical data
- Batch processing
- Use: Dashboards, BI tools



☁️ Cloud vs Self-hosted

Cloud-native

- Pay-as-you-go
- Elastic scaling
- Separates compute & storage
- Quick to deploy

Self-hosted

- More control
- Requires maintenance
- Can be cheaper long-term



🌐 Distributed Systems

When needed: scale, fault tolerance

Challenges: latency, consistency, complexity

Tip: Don't go distributed unless you must!



🔒 Privacy & Compliance

Driven by regulations (e.g. GDPR)

Not just legal — it's ethical

Hard to translate rules to code

Key Takeaway:

Architecture is a series of informed trade-offs.
Make them wisely.



What Makes a Good System?

Think beyond features!
We need systems that are:

- Fast (Performance)
- Scalable
- Reliable
- Maintainable

These are called Non-Functional Requirements (NFRs)

Performance

- Measure with **response time** percentiles (not just averages!)
- Track **throughput** (requests/second, etc.)
- Define **SLAs** (Service Level Agreements)
- Example:
95% of API calls finish under 200ms



Key Takeaway:

"Great systems are not just functional — they're flexible, reliable, and built to grow."



Scalability

- Goal: Handle more load without breaking
- Break down tasks into **independent parts**
- Scale horizontally (more machines, services)
- Analogy: Like a pizza shop hiring more chefs to make pizzas faster

Reliability

- Use **fault tolerance**: system keeps working even if parts fail
- Hardware vs. Software faults
- Human errors? → Build tools to **recover gracefully**
- Use **blameless postmortems** to learn from failures
Mistakes happen. Learn, don't blame!



Maintainability

- Support **operations teams**
- Manage **complexity** → Use clean abstractions
- Enable **easy updates** over time
- Build with modular, well-known components

in @ankitrathi

Good Data Systems

Meeting Non-functional Requirements

A Tour of Data Models

With Query Languages

in @ankitrathi

Relational Model

Used in analytics & data warehousing
Star & Snowflake schemas
Strong SQL support



Structured, reliable, and still going strong after 50+ years.



Graph Model

Highly connected data
Supports recursive queries
Used with Cypher, SPARQL, Datalog



When everything is related to everything



Dataframes

Wide tables with many columns
Perfect for ML & statistical computing
Bridge between databases and machine learning



Emulating Models

Models can mimic each other
But: may get awkward (e.g., recursion in SQL)
Possible but not always elegant



Event Sourcing & CQRS

Append-only log of events
Materialized views for querying
Write-fast, read-smart pattern

Specialized Models

GenBank: for DNA string search
Ledgers: double-entry accounting
Blockchains: distributed ledgers
Full-text search & vector search

Niche needs, custom tools



Polyglot Databases

Databases expanding across models
Relational → JSON
Document → Joins
SQL → Graph

The lines are blurring



Schema: Explicit vs Implicit

- Relational: schema on write
 - Document: schema on read
- Flexibility vs control

Key Takeaway:
Different models for different minds
The key is choosing what fits best

💡 OLTP vs OLAP

OLTP = Fast small reads/writes

- ◆ Indexed access (Primary / Secondary)
- ◆ Handles user-facing apps
- ◆ Uses B-Trees

💡 e.g. Shopping cart updates



OLAP = Complex analytical queries

- ◆ Scans large data volumes
- ◆ Columnar storage + compression
- ◆ Used in reporting, BI tools

💡 e.g. Sales dashboard analytics



Storage & Retrieval

How Databases work?

in @ankitrathi



📁 Storage Engines

Log-Structured Storage

- ✓ Append-only
- ✓ High write throughput
- brick Examples: LSM Trees, SSTables, RocksDB
- blue square Deletes handled by compaction

Update-in-Place Storage

- ✓ Overwrites fixed-size pages
- ✓ Better read latency
- blue square B-Trees dominate relational DBs
- 👉 Use case decides the engine choice!

📘 Index Types

Multidimensional Indexes (R-Trees)

For spatial queries (e.g., lat-long)

blue square "Find all stores near me"

Full-Text Search Index

- 🔍 Finds docs by keywords
- blue square Used in blogs, articles

Vector Search (Semantic)

- 🤖 For similarity search (AI/ML)
- brain "Compares high-dimensional vectors
- "Find docs like this one"



💡 Developer Wisdom

- 🔧 Understand storage = Better tuning
- 🔧 Choose DB based on read/write pattern
- blue square Now you can decode DB docs with confidence!



1 Why Encoding Matters

Converts data structures → bytes (for network or disk)
Influences performance, compatibility, and application design

Key for evolvability (changing systems safely over time)

0 1 0
1 0 1
0 1 0



2 Rolling Upgrades Need Compatibility

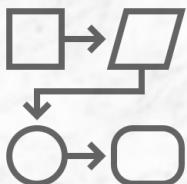
Deploy new versions gradually (rolling upgrades).

Old and new systems must understand each other's data
Goal:

- Backward compatibility (new reads old)
- Forward compatibility (old reads new)

3 Encoding Types & Trade-offs

Type	Pros	Cons
abc Language-Specific	Easy to implement	Tied to language, poor compatibility
text (JSON, XML, CSV)	Human-readable, flexible schema	Loose types, verbose
📦 Binary (Avro, ProtoBuf)	Compact, strict schema, good for APIs	Not human-readable



4 Where Encoding Matters (Dataflow)

-  Database: Encode → store → decode
-  APIs: Encode request → decode/encode response
-  Events: Messages sent between services

Final Nugget

Design for compatibility from Day 1 and make evolution painless

Encoding & Evolution

How data flows?

in @ankitrathi

Why Replicate?

- 💡 **High Availability** - Survive machine/region failures
- 🔌 **Disconnected Operation** - Keep apps running offline
- ⚡ **Low Latency** - Keep data close to users
- 📈 **Scalability** - Handle more reads



Challenges

- 🧠 Not just copying — must handle:
- ❗ Node failures
- 🌐 Network issues
- 🐞 Bugs & data corruption
- ⌚ Replication lag

Replication

Making Data Available Everywhere



Replication Types

in @ankitrathi

Single-Leader

Writes go to one leader
Followers replicate
Easy, but stale reads possible

Multi-Leader

Multiple leaders accept writes
Conflicts need resolution
Good for multi-region writes

Leaderless

Write to multiple nodes
Read from many nodes
Needs conflict detection

4. Consistency Models

- 📌 **Read-After-Write** - See your own write
- 📌 **Monotonic Reads** - No time-travel
- 📌 **Consistent Prefix** - Replies don't come before questions



5. Conflict Resolution

- 🕒 Version vectors - Track concurrent writes
- 🏆 Last Write Wins - Easy, but risky
- 🛠 Manual resolution - Needs dev input
- 🧠 CRDTs - Auto-merge friendly data types

Closing Note

Replication isn't just copying. It's designing for resilience, performance, and peace of mind

Sharding

Scaling Data Across Machines

in @ankitrathi

Why Sharding?

- Too much data for one machine
- Distribute data + query load
- Avoid hot spots (overloaded nodes)
- Enables horizontal scalability



Two Main Sharding Strategies

Key Range Sharding

- Data sorted by key
- Each shard owns a range
- Pros: Fast range queries
- Cons: Risk of hot spots
- Rebalance: Split large ranges



Hash Sharding

- Apply $\text{hash}(\text{key})$
- Shards own hash ranges
- Pros: Even load
- Cons: No efficient range queries
- Rebalance: Move entire shards



Composite Keys

- Use prefix of key to pick shard
- Use suffix for sorting within shard
- Still supports range queries per prefix



Secondary Indexes

Local Secondary Index

- Stored in same shard as data
- Efficient writes
- Reads hit all shards



Global Secondary Index

- Indexed values are sharded separately
- Reads from one shard
- Writes touch many shards



Query Routing

- Coordinator keeps shard-to-node map
- Routes queries to correct shard



Challenges

- Shards work independently
- But multi-shard writes can fail
- What if one succeeds & another doesn't?
- Solution: Covered in next chapter



Key Takeaway

Sharding lets us scale data systems...
...but choosing the right strategy is key to performance, flexibility, and fault tolerance.

Follow:



@ankitrathi



Ankit Rathi (He/Him)

I simplify Data & AI through Visual Storytelling – One Concept at a Time | All views are my own

Noida, Uttar Pradesh, India · [Contact info](#)

[Check My Visual Notes](#) ↗



Bookmark:

[Access all Live Notes here](#)

<https://github.com/ankit-rathi/The-Data-AI-Visual-Notebook>

40+ Visual Notes for **FREE**